Online Recommendation via Particle Thompson Sampling

Michael Alvarino [1]   **Bharat Srikishan** [1]   **Colby Wise** [1]

Abstract

Within the field of recommender systems various forms of collaborative filtering are often used to estimate how users will rate items. One of the most popular methods used when contextual information is not available is Probabilistic Matrix Factorization (PMF). This core method is capable of scaling with a large number of observations and performs well even when restricted to sparse datasets. Unfortunately PMF is limited to offline predictions for a fixed set of users and items.

We present a fast online recommendation system using PMF, Thompson sampling, and particle filtering to provide cold start movie recommendations to users. We also examine the effect of different particle sampling methods on degeneracy of the algorithm.

1. Introduction

PMF is the most common collaborative filtering method for recommendation systems. While it does provide useful recommendations for users it also requires a-priori knowledge of the set of users and items. Real world use cases must provide recommendations for users that have no previous history (cold start users).

Online learning methods are designed to accept new users or items over time without sacrificing the efficiency or accuracy of the algorithm. In Efficient Thompson Sampling for Online Matrix Factorization Recommendation (Kawale et al., 2015) (ETSOMF) an online algorithm is proposed, but is subject to degeneracy, a common problem in particle filtering methods.

A particle filtering algorithm is said be degenerate when, after only a few iterations, all but one particle has negligible weights. We propose several different particle resampling methods which combat the degeneracy problem and compare their performance on the MovieLens 100k Dataset.

2. Model

2.1. Probabilistic Matrix Factorization

Offline PMF computes a low rank factorization of a ratings matrix R where $R = UV^\top$ for $U \in \mathbb{R}^{n \times k}$ and $V \in \mathbb{R}^{m \times k}$. U and V are user and item latent matrices and K is small, for most of our experiments K was less than 10.

The PMF model assigns a Normal prior with zero mean and K dimensional user-specific variance, $\sigma_u^2 I_K$, to the user latent matrix U and a similar distribution to the item latent matrix V. Each rating $r_{i,j}$ is then assigned a Normal distribution centered on $U_i V_j^T$ with rating variance σ^2.

$$U_i \sim \mathcal{N}(0, \sigma_u^2 I_K)$$
$$V_i \sim \mathcal{N}(0, \sigma_v^2 I_K)$$
$$r_{ij}|U, V \sim \mathcal{N}(U_i^\top V_j, \sigma^2)$$

This linear Gaussian model has an analytically solvable conditional posterior, but in order to support online inference we use a particle filtering method as described in ETSOMF.

2.2. Particle Filtering

Particle Filtering (also known as Sequential Monte Carlo) methods assume the state of a system to be unknown and determined by some hidden variables. We would then like to estimate the state of this system using noisy observations and maintaining N particles, or state estimates.

Let $\{X_{0:k}^i, w_k^i\}_1^N$ denote N random measures representative of the posterior of a system, $p(x_{0:k}|z_{1:k})$. Then, we can approximate the true posterior of this system by the following

$$p(x_{0:k}|z_{1:k}) \approx \Sigma_{i=1}^N w_k^i \delta(x_{0:k} - x_{0:k}^i)$$
$$where$$
$$\Sigma_i w_k^i = 1$$

Then, using sequential importance sampling, we can derive update equations for the weight associated with each particle as described in (Arulampalam et al., 2002).

2.3. Particle Thompson Sampling

We take advantage of the graphical structure of our probabilistic model to derive an efficient Rao-Blackwellized particle filter which maintains an estimate of the posterior over time. Additionally, we use Thompson sampling to choose particles and balance exploration with exploitation of our current posterior estimate.

Each particle stores parameters U, V, σ_U, σ_V and we for each update sample $U_{i_t} | V, \sigma_U$ followed by $V_{j_t} | U, \sigma_v$:

$$P(U_i | V, R^o, \sigma, \sigma_U) = P(U_i | V_{rts(i)}, R^o_{i,rts(i)}, \sigma_U, \sigma)$$
$$= \mathcal{N}(U_i | \mu_i^u, (\Lambda_i^u)^{-1})$$
$$\text{where } \mu_i^u = \frac{1}{\sigma^2} (\Lambda_i^u)^{-1} \zeta_i^u$$
$$\Lambda_i^u = \frac{1}{\sigma^2} \sum_{j \in rts(i)} V_j V_j^\top + \frac{1}{\sigma_u^2} I_K$$
$$\zeta_i^u = \sum_{j \in rts(i)} r_{ij}^o V_j$$

Here R^o are the observed ratings and $rts(i)$ is the set of items rated by user i. The update for our sample of V_j mirrors this update.

2.4. Resampling

The original authors of ETSOMF saw degeneracy as a significant problem while training their algorithm. Our primary contribution is the analysis of new resampling methods to combat degeneracy by maximizing the number of unique particles over every iteration of the algorithm. We examin three primary resampling techniques, multinomial, stratified, and systematic resampling.

2.4.1. MULTINOMIAL RESAMPLING

In multinomial resampling the particles are sampled from a multinomial distribution whose parameters are directly defined by the normalized, updated, weights of each particle (Douc & Cappé, 2005).

2.4.2. STRATIFIED RESAMPLING

Stratified resampling breaks the interval into n disjoint sets $(0, 1/n], ..., (\{n-1\}/n, 1]$ and draws independently from the uniform distribution $U((\{i-1\}/n, i/n])$ where i is the i-ith sample (Douc & Cappé, 2005).

2.4.3. SYSTEMATIC RESAMPLING

Systematic resampling deterministically links variables drawn in sub-intervals during stratified sampling. It does so by setting $U^i = (i-1)/n + U$, where U is drawn from $U((0, 1/n])$.

Us generated via systematic resampling are known within the particle sampling literature as empirically good and computationally simple (Douc & Cappé, 2005)

3. Experiments

3.1. Data

We ran our experiments on the MovieLens 100k dataset. We only used the rating data for our online recommendations, however we used the genre data in our evaluation of user preference drift. We subtracted the mean rating from the data to center ratings at 0.

This dataset has 943 users, 1683 movies, and 100,000 ratings.

4. Results

4.1. MSE

Figure 1. Training set mean squared error over time for K = 3 and 10 particles.

We first examine mean squared error on the training set. As we see in Figure 1, it decreases over time as expected. The final MSE in this case is 1.721.

Table 1. Mean squared error statistics on the test set for various choices of K (latent dimensionality) and P (number of particles)

(K, P)	MEAN	STD DEV	MIN	MAX
(2, 2)	2.924	0.355	2.512	3.585
(2, 5)	2.981	0.387	2.552	3.711
(3, 5)	2.540	0.182	2.408	2.900
(3, 10)	2.838	0.507	2.326	3.639
(5, 2)	2.329	0.086	2.199	2.441
(5, 10)	2.367	0.088	2.251	2.519
(5, 20)	2.357	0.062	2.250	2.413

4.2. Cumulative Take Rate

4.3. Hyperparameter Cross Validation

4.4. User Preference Drift

5. Conclusion

References

Arulampalam, M Sanjeev, Maskell, Simon, Gordon, Neil, and Clapp, Tim. A tutorial on particle filters for online nonlinear/non-gaussian bayesian tracking. *IEEE Transactions on signal processing*, 50(2):174–188, 2002.

Cherkassky, Michael and Bornn, Luke. Sequential monte carlo bandits. *arXiv preprint arXiv:1310.1404*, 2013.

Douc, Randal and Cappé, Olivier. Comparison of resampling schemes for particle filtering. In *Image and Signal Processing and Analysis, 2005. ISPA 2005. Proceedings of the 4th International Symposium on*, pp. 64–69. IEEE, 2005.

Kawale, Jaya, Bui, Hung H, Kveton, Branislav, Tran-Thanh, Long, and Chawla, Sanjay. Efficient thompson sampling for online matrix-factorization recommendation. In *Advances in neural information processing systems*, pp. 1297–1305, 2015.

Wang, Qing, Zeng, Chunqiu, Zhou, Wubai, Li, Tao, Shwartz, Larisa, and Grabarnik, Genady Ya. Online interactive collaborative filtering using multi-armed bandit with dependent arms. *arXiv preprint arXiv:1708.03058*, 2017.

Zhao, Xiaoxue, Zhang, Weinan, and Wang, Jun. Interactive collaborative filtering. In *Proceedings of the 22nd ACM international conference on Conference on information & knowledge management*, pp. 1411–1420. ACM, 2013.