

---

# Online Recommendation via Particle Thompson Sampling

---

Michael Alvarino<sup>1</sup> Bharat Srikishan<sup>1</sup> Colby Wise<sup>1</sup>

## Abstract

Online recommendation is a difficult problem with many practical applications. A variety of techniques exist to provide recommendations to users, but many require batch processing and performance tradeoffs. We present a fast online recommendation system using Matrix factorization, Thompson sampling, and particle filtering to provide cold start movie recommendations to users. We also examine user preference drift for different genres over time.

## 1. Introduction

Matrix factorization (MF) is a common tool for recommendation systems. While it provides useful recommendations for users it requires a-priori knowledge of the set of users and items being recommended. Real world use cases have online systems which must provide recommendations for users that have no previous history (cold start users). Another problem in online systems is the problem of user preference drift, as time progresses some users' preferences changes.

We combine the ideas of matrix factorization with bandit algorithms (Zhao et al., 2013) to build a model that solves these problems. Additionally, we use particle filtering to approximate the posterior distribution.

## 2. Model

### 2.1. Probabilistic Matrix Factorization

We compute a low rank factorization of a ratings matrix  $R$  where  $R = UV^\top$ .  $U \in \mathbb{R}^{n \times k}$  and  $V \in \mathbb{R}^{m \times k}$  are user and item latent matrices and  $K$  is small. The probabilistic model is:

$$\begin{aligned} U_i &\sim \mathcal{N}(0, \sigma_u^2 I_K) \\ V_i &\sim \mathcal{N}(0, \sigma_v^2 I_K) \\ r_{ij}|U, V &\sim \mathcal{N}(U_i^\top V_j, \sigma^2) \end{aligned}$$

While this is a linear Gaussian model which has an analytically solvable conditional posterior, we use a sequential

Monte Carlo method to support our online inference. Specifically, we use a particle filter based approach to approximate our posterior.

### 2.2. Particle Thompson Sampling

We use Thompson sampling while running our online linear bandit. This requires incremental updates of the posterior  $U$  and  $V$ . We implement a Rao-Blackwellized particle filter (RBPF) from (Kawale et al., 2015) based on the structure of our model. While iterating through the ratings data for user  $i$  and item  $j$ , we know that other user vectors  $U_{-i}$  are independent of this rating. Also, we know that other item vectors  $V_{-j}$  are also unaffected. This observation allows us to sample only  $U_i$  and  $V_j$ , resulting in an efficient algorithm.

Each particle stores parameters  $U, V, \sigma_U, \sigma_V$  and we sample  $U_{i_t}|V, \sigma_U$  followed by  $V_{j_t}|U, \sigma_V$ :

$$\begin{aligned} P(U_i|V, R^o, \sigma, \sigma_U) &= P(U_i|V_{rts(i)}, R_{i,rts(i)}^o, \sigma_U, \sigma) \\ &= \mathcal{N}(U_i|\mu_i^u, (\Lambda_i^u)^{-1}) \\ \text{where } \mu_i^u &= \frac{1}{\sigma^2} (\Lambda_i^u)^{-1} \zeta_i^u \\ \Lambda_i^u &= \frac{1}{\sigma^2} \sum_{j \in rts(i)} V_j V_j^\top + \frac{1}{\sigma_u^2} I_K \\ \zeta_i^u &= \sum_{j \in rts(i)} r_{ij}^o V_j \end{aligned}$$

Here  $R^o$  are the observed ratings and  $rts(i)$  is the set of items rated by user  $i$ .

## 3. Experiments

### 3.1. Data

We ran our experiments on the MovieLens 100k dataset. We only used the rating data for our online recommendations, however we used the genre data in our evaluation of user preference drift. We subtracted the mean rating from the data to center ratings at 0.

This dataset has 943 users, 1683 movies, and 100,000 ratings.

## 4. Results

### 4.1. RMSE




Figure 1. Training set mean squared error over time for  $K = 3$  and 10 particles.

We first examine mean squared error on the training set. As we see in Figure 1, it decreases over time as expected. The final MSE in this case is 1.721.

Table 1. Mean squared error statistics on the test set for various choices of  $K$  (latent dimensionality) and  $P$  (number of particles)

(K, P)	MEAN	STD DEV	MIN	MAX
(2, 2)	2.924	0.355	2.512	3.585
(2, 5)	2.981	0.387	2.552	3.711
(3, 5)	2.540	0.182	2.408	2.900
(3, 10)	2.838	0.507	2.326	3.639
(5, 2)	2.329	0.086	2.199	2.441
(5, 10)	2.367	0.088	2.251	2.519
(5, 20)	2.357	0.062	2.250	2.413

### 4.2. Cumulative Take Rate

### 4.3. Hyperparameter Cross Validation

### 4.4. User Preference Drift

## 5. Conclusion

## References

- Cherkassky, Michael and Bornn, Luke. Sequential monte carlo bandits. *arXiv preprint arXiv:1310.1404*, 2013.
- Kawale, Jaya, Bui, Hung H, Kveton, Branislav, Tran-Thanh, Long, and Chawla, Sanjay. Efficient thompson sampling for online matrix-factorization recommendation. In *Advances in neural information processing systems*, pp. 1297–1305, 2015.
- Wang, Qing, Zeng, Chunqiu, Zhou, Wubai, Li, Tao, Shwartz,

Larisa, and Grabarnik, Genady Ya. Online interactive collaborative filtering using multi-armed bandit with dependent arms. *arXiv preprint arXiv:1708.03058*, 2017.

Zhao, Xiaoxue, Zhang, Weinan, and Wang, Jun. Interactive collaborative filtering. In *Proceedings of the 22nd ACM international conference on Conference on information & knowledge management*, pp. 1411–1420. ACM, 2013.