

---

# Online Recommendation via Particle Thompson Sampling

---

Michael Alvarino<sup>1</sup> Bharat Srikishan<sup>1</sup> Colby Wise<sup>1</sup>

## Abstract

Within the field of recommender systems various forms of collaborative filtering are often used to estimate how users will rate items. One of the most popular methods used when contextual information is not available is Probabilistic Matrix Factorization (PMF). This core method is capable of scaling with a large number of observations and performs well even when restricted to sparse datasets. Unfortunately PMF is limited to offline predictions for a fixed set of users and items.

We present a fast online recommendation system using PMF, Thompson sampling, and particle filtering to provide cold start movie recommendations to users. We also examine the effect of different particle sampling methods on degeneracy of the algorithm.

## 1. Introduction

PMF is the most common collaborative filtering method for recommendation systems. While it does provide useful recommendations for users it also requires a-priori knowledge of the set of users and items. Real world use cases must provide recommendations for users that have no previous history (cold start users).

Online learning methods are designed to accept new users or items over time without sacrificing the efficiency or accuracy of the algorithm. In Efficient Thompson Sampling for Online Matrix Factorization Recommendation (Kawale et al., 2015) (ETSOMF) an online algorithm is proposed, but is subject to degeneracy, a common problem in particle filtering methods.

A particle filtering algorithm is said to be degenerate when, after only a few iterations, all but one particle has negligible weights. We propose several different particle resampling methods which combat the degeneracy problem and compare their performance on the MovieLens 100k Dataset.

## 2. Model

### 2.1. Probabilistic Matrix Factorization

Offline PMF computes a low rank factorization of a ratings matrix  $R$  where  $R = UV^T$  for  $U \in \mathbb{R}^{n \times k}$  and  $V \in \mathbb{R}^{m \times k}$ .  $U$  and  $V$  are user and item latent matrices and  $K$  is small, for most of our experiments  $K$  was less than 10.

The PMF model assigns a Normal prior with zero mean and  $K$  dimensional user-specific variance,  $\sigma_u^2 I_K$ , to the user latent matrix  $U$  and a similar distribution to the item latent matrix  $V$ . Each rating  $r_{i,j}$  is then assigned a Normal distribution centered on  $U_i V_j^T$  with rating variance  $\sigma^2$ .

$$\begin{aligned} U_i &\sim \mathcal{N}(0, \sigma_u^2 I_K) \\ V_i &\sim \mathcal{N}(0, \sigma_v^2 I_K) \\ r_{ij}|U, V &\sim \mathcal{N}(U_i^T V_j, \sigma^2) \end{aligned}$$

This linear Gaussian model has an analytically solvable conditional posterior, but in order to support online inference we use a particle filtering method as described in ETSOMF.

### 2.2. Particle Filtering

Particle Filtering (also known as Sequential Monte Carlo) methods assume the state of a system to be unknown and determined by some hidden variables. We would then like to estimate the state of this system using noisy observations and maintaining  $N$  particles, or state estimates.

Let  $\{X_{0:k}^i, w_k^i\}_{i=1}^N$  denote  $N$  random measures representative of the posterior of a system,  $p(x_{0:k}|z_{1:k})$ . Then, we can approximate the true posterior of this system by the following

$$\begin{aligned} p(x_{0:k}|z_{1:k}) &\approx \sum_{i=1}^N w_k^i \delta(x_{0:k} - x_{0:k}^i) \\ &\text{where} \\ \sum_i w_k^i &= 1 \end{aligned}$$

Then, using sequential importance sampling, we can derive update equations for the weight associated with each particle as described in (Arulampalam et al., 2002).

### 2.3. Particle Thompson Sampling

We use Thompson sampling while running our online linear bandit. This requires incremental updates of the posterior  $U$  and  $V$ . We implement a Rao-Blackwellized particle filter (RBPF) from (Kawale et al., 2015) based on the structure of our model. While iterating through the ratings data for user  $i$  and item  $j$ , we know that other user vectors  $U_{-i}$  are independent of this rating. Also, we know that other item vectors  $V_{-j}$  are also unaffected. This observation allows us to sample only  $U_i$  and  $V_j$ , resulting in an efficient algorithm.

Each particle stores parameters  $U, V, \sigma_U, \sigma_V$  and we sample  $U_{i_t}|V, \sigma_U$  followed by  $V_{j_t}|U, \sigma_v$ :

$$\begin{aligned} P(U_i|V, R^o, \sigma, \sigma_U) &= P(U_i|V_{rts(i)}, R_{i,rts(i)}^o, \sigma_U, \sigma) \\ &= \mathcal{N}(U_i|\mu_i^u, (\Lambda_i^u)^{-1}) \\ \text{where } \mu_i^u &= \frac{1}{\sigma^2} (\Lambda_i^u)^{-1} \zeta_i^u \\ \Lambda_i^u &= \frac{1}{\sigma^2} \sum_{j \in rts(i)} V_j V_j^\top + \frac{1}{\sigma_u^2} I_K \\ \zeta_i^u &= \sum_{j \in rts(i)} r_{ij}^o V_j \end{aligned}$$

Here  $R^o$  are the observed ratings and  $rts(i)$  is the set of items rated by user  $i$ .

## 3. Experiments

### 3.1. Data

We ran our experiments on the MovieLens 100k dataset. We only used the rating data for our online recommendations, however we used the genre data in our evaluation of user preference drift. We subtracted the mean rating from the data to center ratings at 0.

This dataset has 943 users, 1683 movies, and 100,000 ratings.

## 4. Results

### 4.1. RMSE

We first examine mean squared error on the training set. As we see in Figure 1, it decreases over time as expected. The final MSE in this case is 1.721.

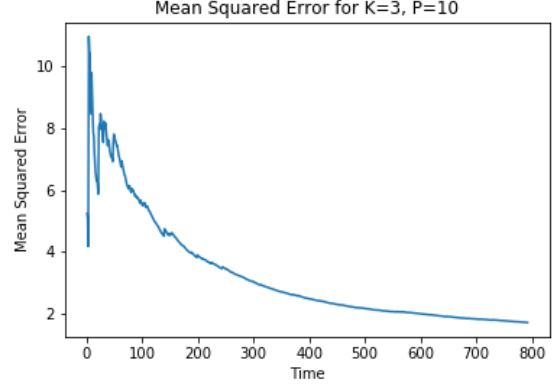


Figure 1. Training set mean squared error over time for  $K = 3$  and 10 particles.

Table 1. Mean squared error statistics on the test set for various choices of  $K$  (latent dimensionality) and  $P$  (number of particles)

(K, P)	MEAN	STD DEV	MIN	MAX
(2, 2)	2.924	0.355	2.512	3.585
(2, 5)	2.981	0.387	2.552	3.711
(3, 5)	2.540	0.182	2.408	2.900
(3, 10)	2.838	0.507	2.326	3.639
(5, 2)	2.329	0.086	2.199	2.441
(5, 10)	2.367	0.088	2.251	2.519
(5, 20)	2.357	0.062	2.250	2.413

### 4.2. Cumulative Take Rate

### 4.3. Hyperparameter Cross Validation

### 4.4. User Preference Drift

## 5. Conclusion

## References

- Arulampalam, M Sanjeev, Maskell, Simon, Gordon, Neil, and Clapp, Tim. A tutorial on particle filters for online nonlinear/non-gaussian bayesian tracking. *IEEE Transactions on signal processing*, 50(2):174–188, 2002.
- Cherkassky, Michael and Bornn, Luke. Sequential monte carlo bandits. *arXiv preprint arXiv:1310.1404*, 2013.
- Kawale, Jaya, Bui, Hung H, Kveton, Branislav, Tran-Thanh, Long, and Chawla, Sanjay. Efficient thompson sampling for online matrix-factorization recommendation. In *Advances in neural information processing systems*, pp. 1297–1305, 2015.
- Wang, Qing, Zeng, Chunqiu, Zhou, Wubai, Li, Tao, Shwartz, Larisa, and Grabarnik, Genady Ya. Online interactive collaborative filtering using multi-armed bandit with dependent arms. *arXiv preprint arXiv:1708.03058*, 2017.

Zhao, Xiaoxue, Zhang, Weinan, and Wang, Jun. Interactive collaborative filtering. In *Proceedings of the 22nd ACM international conference on Conference on information & knowledge management*, pp. 1411–1420. ACM, 2013.